

Jingyu Zhou

EDUCATION

Artificial Intelligence, SJTU

GPA 4.04/4.30 (92.0/100), Rank 4/95

Zhiyuan Honor Program, SJTU

Top 10% of SJTU with advanced coursework

Sept. 2023 – Present

Minhang, Shanghai

Sept. 2023 – Present

Minhang, Shanghai

- **Math:** Discrete Mathematics, Mathematical Analysis, Linear Algebra, Convex Optimization, Probability
- **CS:** Program Design, Data Structure and Algorithm, Machine Learning, Deep Learning

EXPERIENCE

IPADS, SJTU

Research Intern

Dec. 2024 - Present

Minhang, Shanghai

- Implemented KUNSERVE, a parameter-centric approach to optimize GPU memory usage in large language model (LLM) serving. (Codes coming soon)
- Introduced remote attention mechanism via pipeline parallelism, enabling memory borrowing from remote GPUs during load spikes. Thus tackled challenges such as KVCache exchange, balancing memory requirements, and minimizing execution overhead.
- Achieved up to 27.3x reduction in tail TTFT compared to state-of-the-art methods.

RHOS, SJTU & Qing Yuan Research Institute

Research Intern

Sept. 2024 – Jan. 2025

Minhang, Shanghai

- Proposed that the current definition of modalities is inadequate, searching for new definition of them.
- Joined a research team focused on multimodal machine learning, with an emphasis on Large Language Models (LLMs).

X-LANCE, SJTU

Research Intern

Feb. 2024 – Sept. 2024

Minhang, Shanghai

- Contributed to the preparation of a tutorial report on *TTS Based on Discrete Representations* for NCMMSC 2024, specifically focusing on ASR probe experiments.
- Conducted a comprehensive comparison of three types of TTS models (reconstruction, classification, and hybrid) against traditional continuous representations, demonstrating that discrete representations are comparable or superior performance to continuous representations under specific conditions. This finding offers valuable insights for future TTS model development.

HONORS AND AWARDS

Shanghai Scholarship, 2024

- Awarded to the **TOP 0.02%** of students in Shanghai.

Zhiyuan Honors Scholarship, SJTU, 2023 & 2024

- Awarded to the **TOP 5%** of students in the Zhiyuan Honor Program.

Merit Scholarship, B level, SJTU, 2024

- Awarded to the **TOP 5%** of students at SJTU.

PROJECTS

KunServe

- Contributed to the framework of KunServe, including LLaMa backend and LLM scheduling part.
- Implemented pipeline and tensor parallelism coordinated with remote attention mechanism.

SKILLS & INTERESTS

Language: Mandarin (native), English (TOEFL under preparation)

Programming: C/Cpp, Python, Pytorch

Interests: Large Language Model, Machine Learning System

jy-zhou@sjtu.edu.cn | [Homepage](#)